

ECON 3640–001

Problem Set 3 – Solutions

Marcio Santetti

Spring 2022

Problem 1

(a) Simulate 1,000 values drawn from a **Poisson**-distributed random variable, with rate parameter 16. Do not forget to set a **seed** for reproducibility.

ANSWER:

```
set.seed(123)
x_a <- rpois(n = 1000, lambda = 16)
```

(b) Generate 10,000 values drawn from a **Beta**-distributed random variable, with shape parameters 10 and 15. Do not forget to set a **seed** for reproducibility.

ANSWER:

```
set.seed(123)
x_b <- rbeta(n = 10000, shape1 = 10, shape2 = 15)
```

(c) Generate 1,000 values drawn from a **uniformly**-distributed random variable, ranging from 0 to 20. Do not forget to set a **seed** for reproducibility.

ANSWER:


See part (d).

(d) Repeat part (c), now simulating 10,000 values. Produce a histogram for both parts, and compare your results.

ANSWER:

```
set.seed(123)
x_c <- runif(n = 1000, min = 0, max = 20)


x_c %>%
  as_tibble() %>%
  ggplot(aes(x = value)) +
  geom_histogram(color = "white", alpha = 0.7, fill = "pink") +
  labs(title = "n = 10,000 draws") +
  labs(title = "n = 1,000 draws")
```



```
# Notice the difference on the vertical axis!

set.seed(123)
x_d <- runif(n = 10000, min = 0, max = 20)

x_d %>%
  as_tibble() %>%
  ggplot(aes(x = value)) +
  geom_histogram(color = "white", alpha = 0.7) +
  labs(title = "n = 10,000 draws")
```




(e) Generate 5,000 draws from a Normal distribution, with mean 5 and standard deviation 2. Generate these simulated data's density curve, and put, on the same plot, its long-run probability density function (PDF) curve. Compare your results. Do not forget to set a **seed** for reproducibility.

ANSWER:

```
set.seed(123)
x_e <- rnorm(n = 5000, mean = 5, sd = 2)

x_e %>%
  as_tibble() %>%
  ggplot(aes(x = value)) +
  geom_density() +
  stat_function(fun = dnorm, args = list(mean = 5, sd = 2),
               size = 1.2, alpha = 0.8, color = "#8b9dc3") +
  stat_function(fun = dnorm, args = list(mean = 5, sd = 2),
               geom = "area", fill = "#8b9dc3", alpha = 0.35)
```



Problem 2

Consider the following model:

$$Y|\theta \sim \text{Binomial}(n, \theta)$$
$$\theta \sim \text{Beta}(3, 8)$$

Approximate the posterior distribution, $\theta|Y$, using **grid approximation**. How "thin" the grid will be is up to you, but make sure to play around with different values.

Lastly, assume that you observe $n = 100$ trials, with $Y = 5$ successes.

ANSWER:

```
# Setting up the grid, prior, and likelihood:


data_grid <- tibble(
  theta_grid = seq(from = 0, to = 1, by = 0.01),
  prior = dbeta(theta_grid, shape1 = 3, shape2 = 8),
  likelihood = dbinom(x = 5, size = 100, prob = theta_grid)
)

# Calculating the unstandardized and standardized posteriors:

data_grid <- data_grid %>%
  mutate(unstd_posterior = prior * likelihood,
         std_posterior = unstd_posterior/sum(unstd_posterior))

# Plotting the posterior:

data_grid %>%
  ggplot(aes(x = theta_grid, y = std_posterior)) +
  geom_line(size = .8)
```



Problem 3

Use **grid approximation** to estimate the posterior distribution for a Gamma-Poisson model, where you observe $n = 3$ independent data points $(Y_1, Y_2, Y_3) = (0, 1, 0)$. The likelihood and prior are the following:


$$Y_i | \lambda \sim \text{Poisson}(\lambda)$$
$$\lambda \sim \text{Gamma}(20, 5)$$

ANSWER:

```
# Setting up the grid, prior, and likelihood:
data_grid_gp <- tibble(
  lambda_grid = seq(from = 0, to = 10, by = 0.01),
  prior = dgamma(x = lambda_grid, shape = 20, rate = 5),
  likelihood = dpois(x = 0, lambda = lambda_grid) *
    dpois(x = 1, lambda = lambda_grid) *
    dpois(x = 0, lambda = lambda_grid) ## recall that we have 3 independent observations.
)
```

```
# Calculating the unstandardized and standardized posteriors:
data_grid_gp <- data_grid_gp %>%
  mutate(unstd_posterior = prior * likelihood,
    std_posterior = unstd_posterior / sum(unstd_posterior))
```

```
# Plotting the posterior:
data_grid_gp %>%
  ggplot(aes(x = lambda_grid, y = std_posterior)) +
  geom_line(size = .8)
```



Problem 4

Using `stan` syntax, set up the `model{}` block of a model string for an MCMC simulation for the following priors and likelihood functions:

(a) $Y|\theta \sim \text{Binomial}(20, \theta)$ and $\theta \sim \text{Beta}(1, 1)$;

ANSWER:

```
model_a <- "  
  data {  
  
    int<lower=0, upper=20> Y;  
  
  }  
  
  parameters {  
  
    real<lower=0, upper=1> theta;  
  
  }  
  
  model {  
  
    Y ~ binomial(20, theta);  
    theta ~ beta(1, 1);  
  
  }  
"
```

(b) $Y|\lambda \sim \text{Poisson}(\lambda)$ and $\lambda \sim \text{Gamma}(4, 2)$;

ANSWER:

```
model_b <- "  
  data {  
  
    int<lower=0> n;  
    int<lower=0, upper=n> Y;  
  
  }  
  
  parameters {  
  
    real<lower=0> lambda;  
  
  }  
  
  model {  
  
    Y ~ poisson(lambda);  
    lambda ~ gamma(4, 2);  
  
  }  
"
```

(c) $Y|\mu \sim \mathcal{N}(\mu, 1)$ and $\mu \sim \mathcal{N}(0, 10)$;

ANSWER:

```
model_c <- "  
  data {  
    real Y;  
  
  }  
  parameters {  
    real mu;  
  
  }  
  model {  
    Y ~ normal(mu, 1);    // Here, sigma is known.  
    mu ~ normal(0, 10);  
  
  }  
"
```

(d) $Y|\mu, \sigma \sim \mathcal{N}(\mu, \sigma)$, with $\mu \sim \mathcal{N}(0, 10)$ and $\sigma \sim \text{Uniform}(0, 100)$.

ANSWER:

```
model_d <- "  
  data {  
    real Y;  
  
  }  
  parameters {  
    real mu;  
    real <lower=0> sigma; // the standard deviation cannot be negative.  
  
  }  
  model {  
    Y ~ normal(mu, sigma); // now, sigma is unknown.  
    mu ~ normal(0, 10);  
    sigma ~ uniform(0, 100);  
  
  }  
"
```


Problem 5

Using the information from **Problem 2**, set up an **MCMC** simulation using `rstan` with 4 chains and 10,000 iterations. Then, do the following:


```
model_5 <- "  
  
  data {  
  
    int<lower=0> n;  
    int<lower=0, upper=n> Y;  
  
  }  
  
  parameters {  
  
    real<lower=0, upper=1> theta;  
  
  }  
  
  model {  
  
    Y ~ binomial(100, theta);  
    theta ~ beta(3, 8);  
  
  }  
"
```

library(rstan)

```
set.seed(123)  ## don't forget to set a seed!  
  
model_5_mcmc <- stan(  
  model_code = model_5,      ## the model from before.  
  data = list(Y = 5, n = 100), ## hyperparameters were defined above.  
  chains = 4, iter = 5000 * 2 ## run 4 parallel Markov chains, with 10,000 simulations.  
)
```


(a) Generate trace plots for your Markov chains. Did they converge?

ANSWER:



(b) Generate the density plots of your parameter. Are these densities similar to each other?

ANSWER:



(c) What is the posterior mean of your parameter?

ANSWER:

```
model_5_mcmc
```

```
# Inference for Stan model: fab6b48f5279803ac7a47101cf196f51.
# 4 chains, each with iter=10000; warmup=5000; thin=1;
# post-warmup draws per chain=5000, total post-warmup draws=20000.

#               mean se_mean   sd  2.5%   25%   50%   75%  97.5% n_eff Rhat
# theta    0.07    0.00 0.02   0.03   0.05   0.07   0.09   0.13  6769   1
# lp__    -29.26    0.01 0.72  -31.29 -29.43 -28.98 -28.80 -28.75  6043   1
```

The posterior mean for θ is 0.07, given our data and prior model.

(d) Interpret the 95% credibility interval you've obtained from your simulation.

ANSWER:

From the summary above, we can conclude that, given our data and prior, about 95% of our posterior observations lie between 0.03 and 0.13.

Problem 6

Using the `howell.csv` data set (available on Canvas), run an MCMC simulation to approximate the posterior distribution for the average *height*, but only for adults (i.e., individuals older than 18 years old). In other words, make sure to apply the `filter()` function to your data set first.

You may assume, for simplicity, that the standard deviation parameter, σ , is constant and known. You may use any value you want, as long as it makes sense. Recall that, if your data are normally distributed, approximately 95% of the observations lie within 2 standard deviations of the mean, or $\mu \pm 2 \cdot \sigma$. Take this detail into account when setting up your Bayesian model.

Answer the same questions from Problem 5.

ANSWER:

```
howell <- read_csv("howell.csv")

howell_adults <- howell %>%
  filter(age > 18) ## just the adult individuals here.
```

If we assume that the average height of 170 cm for an adult population is normally distributed with a standard deviation of 15 cm, approximately 95% of the individuals will have a height between 140 and 200 cm.

Let's use these as our priors.

```

model_6 <- "
  data {
    real<lower=0> Y[346]; // the filtered data set has 346 adults.

  }

  parameters {

    real<lower=0> mu; // we cannot have negative heights.

  }

  model {
    Y ~ normal(mu, 15); // sigma is assumed to be known here.
    mu ~ normal(170, 15);

  }
"

```

Running the MCMC sampler:


```

set.seed(123)


model_6_mcmc <- stan(
  model_code = model_6,
  data = list(Y = howell_adults$height), ## selecting the 'height' column.
  chains = 4,
  iter = 5000 * 2
)

```

(a) Trace plot:



(b) Density plot:



(c) The posterior mean for μ is 154.68 cm, given our data and prior model.

(d)

model_6_mcmc

```
# Inference for Stan model: 6043b9632a6b63cd5fdc40421532ba54.
# 4 chains, each with iter=10000; warmup=5000; thin=1;
# post-warmup draws per chain=5000, total post-warmup draws=20000.

#               mean se_mean   sd  2.5%  25%   50%   75%  97.5% n_eff Rhat
# mu          54.68    0.01 0.81 153.12 154.13 154.67 155.23 156.26 6862   1
# lp__        -42.31    0.01 0.71 -44.32 -42.48 -42.04 -41.86 -41.81 8806   1
```

From the summary above, we can conclude that, given our data and prior, about 95% of our posterior observations for the average height of adults (μ) lies between 153.12 cm and 156.26 cm.

Problem 7

Still using the `howell.csv` data set, now estimate a Bayesian model for the posterior distributions for *weights* (in kilograms). This time, you will have to set a prior for σ as well.

Here, make sure to play around with models including every individual in the sample, but also for adults only.

As a bonus, you may estimate models also based on `gender` since, these data are also available in this data set. With different research questions in mind, make sure to think about your priors and adjust your data accordingly.

ANSWER:

Let's assume that weights of individuals are normally distributed, with mean and standard deviation parameters μ and σ , respectively.

Since we don't know σ , we need to set priors for both parameters.

Priors: $\mu \sim \mathcal{N}(\theta, \tau)$

$\sigma \sim \text{Uniform}(0, 20)$

Let's set up now the model string:

```
model_7 <- "
  data {
    real<lower=0> Y[544]; // using the whole sample.
  }
  parameters {
    real<lower=0> mu;      // we cannot have negative weights.
    real <lower=0, upper=20> sigma;
  }
  model {
    Y ~ normal(mu, sigma);
    mu ~ normal(65, 10);  // assuming ~ N(60,20) for the average weight.
    sigma ~ uniform(0, 20); // assuming the SD to be uniformly distributed between 0 and 20.
  }
"
```

Then, we use `rstan` to run MCMC simulations.

```
set.seed(123)


model_7_mcmc <- stan(
  model_code = model_7,
  data = list(Y = howell$weight),
  chains = 4,
  iter = 5000 * 2
)
```


```
# Inference for Stan model: c80f0e81d8ee9675915a189dd491cfa7.
# 4 chains, each with iter=10000; warmup=5000; thin=1;
# post-warmup draws per chain=5000, total post-warmup draws=20000.
```

#	mean	se_mean	sd	2.5%	25%	50%	75%	97.5%	n_eff	Rhat
# mu	35.73	0.00	0.64	34.47	35.30	35.73	36.17	36.99	16990	1
# sigma	14.75	0.00	0.45	13.90	14.45	14.74	15.05	15.68	17416	1
# lp__	-1734.80	0.01	1.05	-1737.64	-1735.20	-1734.47	-1734.06	-1733.79	8261	1

From this model, we see that, given our data and priors, the posterior average weight is of 35.73 kg, with a posterior mean standard deviation of 14.75 kg.

Let's have a look at trace and density plots:





The MCMC sampler has converged (as seen in the trace plots) and the density plots for all 4 chains are very similar, indicating a well-specified Bayesian posterior model.

Feel free to further customize and modify priors and the sample you will be investigating here, as suggested in the problem.